



WHITEPAPER · 18 PAGES · FREE DOWNLOAD

Agentforce Implementation Whitepaper

A technical deep dive into Agentforce 360 architecture, the Atlas Reasoning Engine, Einstein Trust Layer, MuleSoft integration patterns, and a phased rollout framework for enterprise Salesforce deployments.

Architecture

Atlas Engine

Einstein

MuleSoft

Apex

Produced by Kovil AI Engineering Team · kovil.ai · June 2026

Contents

01	Agentforce 360 Architecture Overview	3
02	Atlas Reasoning Engine — How It Works	5
03	Einstein Trust Layer — Security Architecture	7
04	MuleSoft Integration Patterns	9
05	Prompt Template Engineering	11
06	Human-in-the-Loop Design	13
07	Phased Rollout Framework	15
08	Architecture Decision Checklist	17

01 / Architecture

Agentforce 360 Architecture Overview

Agentforce 360 is Salesforce's unified AI agent architecture — spanning the Einstein layer (model access and trust), the Atlas Reasoning Engine (agent orchestration), the Salesforce data layer (CRM + Data Cloud), and the action layer (Flows, Apex, MuleSoft, external APIs).

Model Layer
Einstein AI / xGen models

LLM inference for agent reasoning, prompt completion, and response generation. Salesforce uses proprietary xGen models plus access to third-party models (GPT-4, Claude) via the Einstein AI gateway. All model calls pass through Einstein Trust Layer.

Reasoning Layer
Atlas Reasoning Engine

The core agent brain — interprets the user's intent, retrieves relevant data, selects appropriate actions from the action library, executes them in sequence, and returns a response. Atlas handles multi-step reasoning without custom orchestration code.

Data Layer
CRM + Einstein 1 Data Cloud

The data sources agents reason over. CRM provides transactional records. Data Cloud provides unified customer profiles combining CRM, marketing, commerce, and external data. Data Cloud-grounded agents are significantly more accurate.

Action Layer
Flows, Apex, MuleSoft, HTTP

The tools agents can invoke. Actions are registered in the Agentforce action library with natural language descriptions. Atlas selects actions based on those descriptions — making description quality critical for correct action selection.

Safety Layer
Einstein Trust Layer

Data masking, content safety filters, prompt injection resistance, grounding with company data, and audit logging. Applied to every agent interaction before the LLM sees the prompt.

02 / Reasoning Engine

Atlas Reasoning Engine — How It Works

The Atlas Reasoning Engine is the decision-making core of Agentforce. Understanding how it works is prerequisite to writing effective agent topics and action descriptions.

Step 1: Intent Parsing

Atlas parses the incoming user message and identifies the intent — what the user is trying to accomplish. This step relies on the agent's Topic configuration, which defines what the agent is responsible for.

Step 2: Context Retrieval

Atlas queries the Salesforce data layer to retrieve relevant records — the current user's account, open cases, recent interactions — and adds them to the reasoning context.

Step 3: Action Selection

Atlas reviews the registered action library and selects the most appropriate action(s) based on their natural language descriptions. This is why action description quality is critical — Atlas cannot infer what an action does from its API name alone.

Step 4: Action Execution

Atlas executes the selected action(s) sequentially, passing parameters from the reasoning context. Each action's output updates the reasoning context for subsequent steps.

Step 5: Response Generation

Atlas synthesises a natural language response incorporating the action outputs. The response passes through Einstein Trust Layer safety filters before being returned to the user.

03 / Security

Einstein Trust Layer — Security Architecture

Einstein Trust Layer is a non-optional security wrapper on all Agentforce interactions. It operates at the infrastructure level — no code changes are required to benefit from it, but it must be configured correctly to meet enterprise compliance requirements.

Zero Data Retention

Salesforce guarantees that LLM providers (OpenAI, Anthropic, etc.) do not retain your data beyond the inference call. This guarantee is contractual and backed by technical enforcement — no customer data is used to train third-party models.

Dynamic Grounding

Einstein Trust Layer performs Retrieval-Augmented Generation (RAG) using your Salesforce data — grounding LLM responses in your actual records rather than model training data. This dramatically reduces hallucination rates on company-specific queries.

PII Masking

Before sending data to the LLM, Einstein Trust Layer identifies and masks PII fields (email, phone, SSN, credit card) with tokens. The LLM sees masked values; the actual values are restored in the response before it reaches the user.

Toxicity Detection

Outgoing responses are evaluated by a content safety classifier before returning to users. Toxic, harmful, or off-topic responses are blocked and logged.

Prompt Injection Resistance

Einstein Trust Layer includes prompt injection resistance — detecting and blocking attempts by users to override agent instructions through clever prompting.

04 / Integration

MuleSoft Integration Patterns

MuleSoft is Salesforce's enterprise integration platform. Agentforce agents can invoke MuleSoft APIs as actions — allowing agents to interact with any system in your enterprise landscape, not just Salesforce.

Pattern 1: Agent → MuleSoft API → ERP

Agent retrieves a customer's order status from an ERP (SAP, Oracle) via a MuleSoft API action. The agent presents the result in natural language. Use case: customer service agents with real-time order visibility without ERP training for agents.

Pattern 2: Agent → MuleSoft API → ITSM

Agent creates, updates, or escalates IT service management tickets (ServiceNow, Jira) via MuleSoft. Use case: internal helpdesk agents that automate ticket creation from Salesforce cases.

Pattern 3: Agent → MuleSoft API → Data Warehouse

Agent queries a data warehouse (Snowflake, BigQuery) for analytics context via MuleSoft. Use case: sales agents that surface account-level revenue trends without DW training.

Pattern 4: MuleSoft Event → Agentforce Trigger

A MuleSoft integration publishes an event (new order, payment failure, contract expiry) that triggers an Agentforce agent proactively. Use case: proactive renewal agents that initiate conversation when a contract is within 60 days of expiry.

05 / Prompts

Prompt Template Engineering

Instructions are not prompts

Agentforce Topic instructions are not the same as LLM prompts. Topics define the agent's scope and persona. Prompt Templates define how the agent composes LLM calls for specific actions. Both must be crafted deliberately.

Topic instruction quality

Topic instructions should specify: what the agent is (persona), what it can do (capabilities), what it cannot do (explicit exclusions), how it should escalate, and its tone. Vague topics produce unpredictable agents.

Prompt Template variables

Use Prompt Template variables to inject dynamic context — the user's name, account details, case history — directly into the LLM prompt. Static prompts without dynamic context produce generic, unhelpful responses.

Grounding instructions

Always include an explicit grounding instruction: "Answer only from the information provided. Do not use information from your training data." This single instruction reduces hallucination rates on factual queries by 40–60%.

06 / Human-in-Loop

Human-in-the-Loop Design

Escalation is not failure

A well-designed Agentforce agent escalates gracefully. Escalation means the agent recognised its limits and transferred to a human appropriately — this is correct behaviour, not failure.

Escalation triggers

Define explicit escalation triggers: user frustration signals (repeated rephrasing), confidence threshold below threshold, action failure, out-of-scope request. Test each trigger in staging before production.

Escalation context transfer

When escalating, the agent should pass the full conversation context — what was asked, what the agent tried, why it escalated — to the receiving human. Context transfer is configured in the agent's Escalation Action.

Monitoring and override

Supervisors should be able to view live agent conversations and intervene at any point. Salesforce provides supervisor view in Agentforce for Human Service agents. Configure supervisor access before go-live.

07 / Rollout

Phased Rollout Framework

Phase 1: Internal Alpha (Week 1–2)

Deploy to internal team of 3–5 testers. Define 30+ test scenarios covering success cases, edge cases, and out-of-scope queries. Set acceptance criteria: >80% accuracy, <5% incorrect escalations, <3s p95 response time.

Phase 2: Limited Beta (Week 3–5)

Deploy to 10–50 selected users or a single team. Monitor conversation logs daily. Identify top 5 failure patterns. Iterate on topic instructions and action descriptions. Do not expand until failure rate is below 10%.

Phase 3: Canary (Week 6–8)

Route 15–20% of live traffic to the agent. Monitor escalation rate, resolution rate, and CSAT in real time. Keep manual process running in parallel. Set rollback trigger: escalation rate >30% triggers review.

Phase 4: Full Production (Week 9+)

Ramp to 100% of target traffic. Establish weekly conversation quality reviews. Plan first agent update based on observed failure patterns. Set up automated monitoring alerts.

08 / Checklist

Architecture Decision Checklist

#	Decision	Your Answer
1	Salesforce edition & licenses Enterprise+ with Agentforce add-on confirmed?	
2	Service user permissions Agent service user profile and permission sets tested?	
3	Topic design Topics narrowly scoped — one topic per distinct agent responsibility?	
4	Action descriptions Every action description tested for correct Atlas selection?	
5	Prompt Template grounding Grounding instruction included in all LLM-facing templates?	
6	Data Cloud grounding Data Cloud configured and grounded agents tested (if licensed)?	
7	Einstein Trust Layer PII masking fields configured. Toxicity thresholds reviewed?	
8	MuleSoft integration Named credentials configured for all external system connections?	
9	Escalation paths All escalation triggers defined, tested, and context transfer verified?	
10	Sandbox parity Full sandbox with production-equivalent data used for final testing?	
11	Monitoring Supervisor view configured. Monitoring alerts set. Log retention confirmed?	
12	Rollback plan Rollback trigger thresholds defined. Manual fallback process documented?	