

THE AGENTFORCE READINESS GUIDE

A complete framework for assessing and closing your Agentforce readiness gaps — from data quality to governance.

by **Kovil AI**

kovil.ai/agentforce

Table of Contents

- 1. Introduction: The Agentforce Readiness Gap**
- 2. Pillar 1: Data Quality & Availability**
- 3. Pillar 2: Org Architecture**
- 4. Pillar 3: Use Case Prioritisation**
- 5. Pillar 4: Team Enablement**
- 6. Pillar 5: Governance & Compliance**
- 7. The Readiness Scorecard**
- 8. Your Next Steps**

Chapter 1: Introduction — The Agentforce Readiness Gap

Agentforce represents the most significant shift in enterprise software since the introduction of cloud CRM. Yet despite widespread licence adoption, the vast majority of Salesforce organisations have zero AI agents running in production. The gap between acquiring access to Agentforce and successfully deploying it is wide — and it is entirely preventable.

At Kovil AI, we have conducted readiness assessments across dozens of organisations in Southeast Asia, Australia, and the Middle East. Our findings consistently reveal five root-cause gaps that prevent organisations from crossing the line from pilot to production.

Key Insight

73% of organisations with Agentforce licences have not deployed a single production agent. The bottleneck is not technology — it is readiness across data, architecture, process, people, and governance.

The Five Readiness Gaps

Gap 1: Data Quality

Agents require clean, structured, accessible data. Most CRM instances contain incomplete records, duplicate accounts, and disconnected data silos that cause agents to produce inaccurate or hallucinated responses.

Gap 2: Org Architecture

Without the right permission sets, named credentials, API configurations, and sandbox strategy, agents cannot access the systems they need to take action.

Gap 3: Use Case Selection

Organisations often start with the wrong use cases — choosing ones that are either too complex technically or deliver too little business value to justify the effort.

Gap 4: Team Capability

Agentforce requires a blend of Salesforce administration, prompt engineering, AI quality assurance, and change management skills that few teams currently possess.

Gap 5: Governance

Regulatory obligations, PII handling requirements, and internal approval workflows must be designed into the agent — not bolted on afterwards.

This guide is structured around these five pillars. Each chapter provides a practical assessment framework, checklist, and concrete guidance. The final chapter consolidates everything into a 25-question scorecard so you can benchmark your organisation's readiness in under 30 minutes.

Chapter 2: Pillar 1 — Data Quality & Availability

An Agentforce agent is only as good as the data it can access. The Einstein Trust Layer ensures that data never leaves Salesforce's secure boundary, but if the underlying data is incomplete, stale, or siloed, the agent will confidently produce wrong answers.

What Data Does an Agent Need?

The answer depends on the use case, but all agents share a common set of data requirements. Service agents need complete case history, product entitlements, and SLA records. Sales agents need accurate account hierarchies, opportunity stages, and contact relationships. Field service agents need asset records, maintenance schedules, and territory assignments.

Data Quality Principle

Garbage in, garbage out applies to AI more severely than any previous technology. An agent that confidently delivers wrong information damages customer trust far more than no agent at all.

Measuring Your Data Readiness Baseline

Before remediating data quality issues, you first need to measure them. Salesforce provides several native tools for this purpose, and the assessment process typically reveals that data quality problems are more concentrated than organisations expect — often 20% of records account for 80% of the quality issues.

Run the following SOQL queries in Salesforce Workbench or Developer Console to establish your baseline metrics. Document these numbers before remediation begins so you can track progress and demonstrate improvement to stakeholders.

Metric	SOQL / Method	Target Threshold	Action if Below Target
Account field completion	COUNT fields with null values per record	>80% complete	Data enrichment sprint
Duplicate account rate	Duplicate Rules report in Setup	<5% duplicates	Merge campaign + Prevention Rule
Contact-Account linkage	Contacts WHERE AccountId = null	<2% orphan contacts	Manual review and link
Opportunity staleness	Opps WHERE LastModifiedDate < 90 days	0% stale active opps	Pipeline hygiene campaign

Metric	SOQL / Method	Target Threshold	Action if Below Target
Case resolution rate	Closed Cases WHERE Resolution = null	<5% missing resolution	Auto-close rule update
Email validity	Contacts WHERE Email invalid format	<1% invalid emails	Email validation flow

Data Enrichment Strategies

Where internal data is incomplete, external enrichment can fill the gaps. Salesforce has native integrations with data enrichment providers including Dun & Bradstreet, ZoomInfo, and Clearbit. However, enrichment is not a substitute for data governance — enriched data must be validated against business rules before being exposed to agents.

- Firmographic enrichment (Industry, Employee Count, Revenue) for Account records using D&B; or ZoomInfo
- Email and phone validation for Contact records using a real-time verification service
- Address standardisation and geocoding for shipping and billing address fields
- Technographic data (which tech stack a prospect uses) for sales agent personalisation
- Social profile linking for Contact records where relevant to the agent use case

Real-Time Data Access Patterns

Some agent use cases require data that changes rapidly — stock levels, booking availability, live ticket queues. For these cases, static CRM data is insufficient and real-time data access patterns are required. Salesforce Data Cloud's real-time event streaming and Calculated Insights provide the mechanism for surfacing live operational data to agents without duplicating it into Salesforce CRM.

The key architectural decision for real-time data is whether to pull data at query time (via an Apex Action or External Service callout in the agent flow) or to push data into Data Cloud in near-real-time (via Salesforce Streaming API, Platform Events, or a Kafka connector). Pull patterns are simpler but add latency; push patterns are lower latency but require more infrastructure.

Clean CRM Data: The Baseline Requirements

Account completeness

All accounts must have: Name, Industry, BillingCountry, Type, and OwnerId populated.

Contact accuracy

Contact records must link to accounts with valid email and phone where applicable.

Duplicate management

Merge or suppress duplicate records — agents cannot resolve ambiguity.

Opportunity hygiene

Stage, CloseDate, and Amount must be current; stale pipe corrupts forecasts.

Case history

Closed cases must have Resolution populated — agents use this for RAG retrieval.

Unified Customer Profiles via Data Cloud

Agentforce performs best when it can access a unified customer profile that aggregates data from CRM, marketing automation, commerce, and external systems. Salesforce Data Cloud provides this unified profile layer through identity resolution and real-time data streams.

Organisations without Data Cloud can still deploy agents, but they will operate with a narrower data window and may require additional custom actions to retrieve cross-system context.

Data Quality Self-Assessment Checklist

	Item	Status
[]	Account records have >80% field completion rate for key fields	_____
[]	Duplicate accounts reduced to <5% of total account volume	_____
[]	Contact-to-Account relationships are accurate and complete	_____
[]	All active Opportunities have current Stage, Amount, and CloseDate	_____
[]	Closed Cases have Resolution and Root Cause populated	_____
[]	Data Cloud or equivalent unified profile exists for customer 360	_____
[]	Real-time data streams are available for time-sensitive use cases	_____
[]	Data dictionary / field glossary is maintained and current	_____
[]	PII fields are identified and tagged in the data model	_____
[]	Data quality dashboard is in place and monitored weekly	_____

Chapter 3: Pillar 2 — Org Architecture

A technically sound Salesforce org is the foundation on which Agentforce runs. Without the right architecture in place, agents will fail silently — unable to access external systems, hitting API limits mid-conversation, or operating without appropriate security controls.

Permission Sets & Security Model

Agentforce agents run in the context of a dedicated Integration or Bot User. This user must have precisely scoped permissions — sufficient to perform its assigned actions, but no broader. Over-permissioned agent users represent a security risk; under-permissioned users cause runtime failures.

Best practice is to create a dedicated Permission Set for each agent type, containing only the object-level and field-level permissions required for its declared actions. Assign this Permission Set to the agent's Integration User exclusively.

Architecture Rule

Never run Agentforce under a System Administrator profile. Create a purpose-built Integration User with the minimum permissions required. Audit these permissions before go-live.

Permission Set Architecture: Step-by-Step

The correct sequence for configuring the Agentforce security model is: first, identify all the objects and fields your agent actions will read or write. Second, create a Permission Set that grants exactly those object permissions and field-level access. Third, create a dedicated Integration User (license type: Salesforce Integration) and assign the Permission Set. Fourth, validate by testing agent actions as this user in a sandbox before promoting to production.

Object/Feature	Permission Required	Read vs. Write	Notes
Case	Object: Read, Create, Edit	Both	Write only if agent creates/updates cases
Account	Object: Read	Read only	Rarely needs write access
Contact	Object: Read	Read only	Write only for contact update use cases
Knowledge Article	Object: Read	Read only	Needed for RAG / knowledge retrieval

Object/Feature	Permission Required	Read vs. Write	Notes
Einstein Bot API	System Permission	N/A	Required for all agent deployments
Named Credentials	Assigned Credential	N/A	One per external system
Flow: Run As	Run Flows permission	N/A	If any Flow-based actions are used

Named Credentials & External Callouts

When agents need to call external APIs — ERP systems, payment gateways, logistics platforms — they must do so via Salesforce Named Credentials. Named Credentials store authentication details securely and prevent hardcoded credentials in Apex or Flow.

Salesforce supports three types of Named Credentials: Legacy Named Credentials (deprecated, do not use for new implementations), Named Credentials with External Credentials (current standard, supporting OAuth 2.0, JWT, custom auth), and Per-User Named Credentials (where each Salesforce user has their own auth token to the external system). For Agentforce, use Named Credentials with External Credentials authenticated as the Integration User.

API Limits & Governance

Agentforce agents consume API calls for every action they execute. In a high-volume deployment, a single agent conversation can trigger 5-20 API calls. Multiply this by concurrent users and you can hit daily API limits rapidly.

Salesforce orgs on Enterprise Edition receive 1,000,000 API calls per 24-hour rolling window. At first this seems generous, but an org with 500 concurrent agent users averaging 10 actions per conversation and 20 conversations per day will consume 100,000 API calls from agents alone — before any other integrations, ETL processes, or connected applications are counted.

Org Edition	Daily API Limit	Recommended Agent Headroom	Monitoring Tool
Essentials	5,000	60% (3,000 reserved)	API Usage Dashboard
Professional	100,000	30% (70,000 reserved)	Event Monitoring API
Enterprise	1,000,000	30% (700,000 reserved)	Event Monitoring + Alerts
Unlimited	5,000,000	20% (4M reserved)	Splunk / Datadog integration

Configuration Area	Requirement	Risk if Missing
Integration User	Dedicated user with min. permissions	Security breach / audit failure
Permission Sets	Scoped per agent type	Over-permissioning or runtime errors
Named Credentials	All external APIs registered	Hardcoded creds / callout failures
API Limit Buffer	>=30% daily headroom reserved	Production outages at peak load
Sandbox Strategy	Full Copy sandbox for UAT	Missed bugs reaching production
Connected App	Dedicated OAuth app for agent	Token conflicts / revocation risk

What a Production-Ready Org Looks Like

- A dedicated Integration User per agent type with a scoped Permission Set
- All external callouts registered as Named Credentials with certificate-based auth
- API usage monitored in real-time with alerts at 70% of daily limit
- A Full Copy sandbox refreshed within 30 days used for agent UAT
- Flow and Apex code coverage >75% with no critical Apex errors in the past 30 days
- Health Check score >80 in the target production org
- Einstein Activity Capture or equivalent audit logging enabled



Chapter 4: Pillar 3 — Use Case Prioritisation

Choosing the right first use case is the single most important decision in an Agentforce deployment. The wrong choice leads to months of wasted effort, a failed pilot, and an organisation that concludes "AI doesn't work for us." The right choice delivers visible ROI in weeks and builds organisational confidence for the next phase.

The 2x2 Prioritisation Matrix

Evaluate every candidate use case against two dimensions: Business Impact (revenue, cost reduction, customer satisfaction) and Implementation Complexity (data readiness, integration depth, change management). The ideal first use case sits in the HIGH IMPACT / LOW COMPLEXITY quadrant.

	LOW Complexity	HIGH Complexity
HIGH Impact	IDEAL STARTING POINT Fast wins, high visibility	STRATEGIC BETS Resource-intensive, plan carefully
LOW Impact	QUICK WINS Good for learning, limited ROI	AVOID High effort, low reward

Scoring Your Use Cases

To apply the 2x2 matrix, score each candidate use case on both dimensions using the following rubric. Sum the scores to produce an overall priority ranking for your backlog.

Dimension	Score 1 (Low)	Score 2 (Medium)	Score 3 (High)
Business Impact	Process efficiency only; internal users; cost savings <\$50K/yr	Mixed internal/external; measurable CSAT impact; savings \$50K-\$500K/yr	Customer-facing; revenue impact; NPS/retention improvement; >\$500K/yr
Implementation Simplicity	Complex: multi-system integration, custom data model, compliance review required	Moderate: 1-2 integrations, standard data model, some governance review	Simple: Salesforce-only actions, clean data available, no regulatory constraints

Selection Rule

Your first Agentforce pilot should score 3 on Business Impact and 3 on Implementation Simplicity. If you cannot find such a use case, choose the highest-scoring available option and invest 2 weeks of data or architecture work to move it into the ideal quadrant.

Top 10 Highest-ROI Agentforce Use Cases

Rank	Use Case	Impact	Complexity	Typical ROI
1	Case Deflection Agent	High	Low	30-50% deflection
2	Order Status Enquiry	High	Low	20-40% AHT reduction
3	Lead Qualification	High	Medium	3x pipeline velocity
4	Renewal & Upsell Coach	High	Medium	15-25% renewal lift
5	Internal HR Policy Agent	Medium	Low	40% HR ticket reduction
6	Field Service Scheduling	High	Medium	20% utilisation gain
7	Quote-to-Cash Assistant	High	High	50% quote cycle reduction
8	Product Knowledge Copilot	Medium	Low	25% support accuracy lift
9	Complaint Resolution Agent	High	Medium	CSAT +12-18 points
10	Onboarding Journey Agent	Medium	Medium	35% completion rate lift

Use Case Deep Dive: Case Deflection Agent

As the highest-ROI, lowest-complexity entry point for most organisations, the Case Deflection Agent deserves a closer examination. This agent intercepts incoming service requests — via web chat, email, or self-service portal — and attempts to resolve them autonomously before a case is created and assigned to a human agent.

The agent retrieves the customer's account and recent case history, searches the Knowledge Base for relevant articles, and presents a resolution. If the customer confirms the issue is resolved, the case is never created. If the customer requires human assistance, the agent creates a pre-populated case and routes it to the appropriate queue with context already attached.

Metric	Pre-Agent Baseline	Post-Agent Target	How Measured
Case deflection rate	0%	30-50%	Cases not created / total enquiries
Average handle time	Baseline (varies)	-20 to -40%	Case closure time in Salesforce
First contact resolution	Baseline (varies)	+15-25%	Cases closed without re-open
Agent satisfaction (staff)	Baseline	+10-20 pts	Internal staff survey (quarterly)
Customer satisfaction	Baseline	+8-15 pts CSAT	Post-interaction survey

Chapter 5: Pillar 4 — Team Enablement

Technology without people is shelf-ware. Agentforce requires a cross-functional team to design, build, test, and continuously improve agents. Most organisations underestimate this requirement and attempt to deliver with a single "Salesforce admin who also does AI."

The Core Team: Four Essential Roles

Salesforce Administrator / Architect

Responsible for org configuration, permission sets, data model changes, Flow automation, and Named Credentials. Must understand the existing org architecture deeply. A junior admin is insufficient — this role requires senior-level experience (minimum 3 years, ideally certified as a Salesforce Architect).

AI/Prompt Engineer

Designs and iterates on system prompts, topic instructions, action descriptions, and guardrails. Must understand how the Atlas Reasoning Engine interprets instructions and how to structure prompts to produce deterministic, reliable agent behaviour.

Business Analyst / Process Owner

Translates business requirements into agent design specifications. Owns the conversation flow mapping, escalation paths, and success metric definition. Must have deep domain knowledge of the target process being automated.

UAT Owner / Quality Lead

Designs and executes the test plan. Responsible for adversarial testing (attempting to break the agent), edge case coverage, and sign-off before go-live. Often a senior business stakeholder or customer experience lead.

The Case Against the "Accidental AI Project Manager"

In many organisations, Agentforce is initially assigned to whoever currently manages the Salesforce org — typically a senior administrator who is already at full capacity delivering declarative automation, user training, and release management. This person cannot deliver a successful Agentforce implementation on top of their existing responsibilities, regardless of how talented they are.

A successful first Agentforce deployment requires dedicated resource allocation. The core team should have at least 50% of their time ring-fenced for the project during the build phase. Part-time project delivery is the leading cause of timeline overruns in our implementation experience.

Skills Gap Analysis

Skill Area	Required Level	Typical Gap	Time to Close
Salesforce Admin (declarative)	Senior (3+ yrs)	Usually present	N/A
Salesforce Apex / API	Intermediate	Often absent	3-6 months
Prompt Engineering	Intermediate	Almost always absent	4-8 weeks
AI Testing & Evaluation	Basic	Always absent	2-4 weeks
Data Cloud configuration	Intermediate	Usually absent	6-12 weeks
Change Management / Training	Senior	Often under-resourced	2-4 weeks

Training & Enablement Requirements

Beyond the core delivery team, a broader group of stakeholders needs training to ensure successful adoption. The training curriculum should be differentiated by role — what a Service Manager needs to understand about Agentforce is very different from what a Salesforce Developer needs to know.

Audience	Topics	Format	Duration
Executive Sponsors	AI strategy, ROI measurement, governance oversight	Executive briefing	2 hours
Business Analysts	Use case design, conversation flow mapping, success metrics	Workshop	1 day
Salesforce Admins	Agent Builder, Topics, Actions, debugging tools	Hands-on lab	2 days
Developers (Apex)	Custom Actions, External Services, testing frameworks	Technical workshop	2 days
End Users (Agents)	When to trust agent output, escalation process, feedback loop	Team training	2 hours
IT/Security	Trust Layer, audit logs, Named Credentials, security model	Briefing + Q&A;	3 hours

Build vs Buy vs Partner

Organisations face a build-buy-partner decision for each gap identified. The right answer depends on time-to-market requirements, strategic importance of the capability, and internal development capacity.

Approach	Best For	Typical Timeline	Key Risks
Build Internally	Orgs with strong SF teams + 12+ month runway	6-18 months	Underestimated complexity
Hire Specialists	Orgs scaling long-term AI programmes	3-6 months to onboard	Talent scarcity, retention
Partner (SI/ISV)	Fast time-to-market, knowledge transfer needed	4-12 weeks	Partner quality variance
Kovil AI Partnership	SE Asia / ANZ orgs needing Agentforce specialists	2-8 weeks	Lowest — fixed-scope delivery



Chapter 6: Pillar 5 — Governance & Compliance

Governance is the pillar most often left until last — and the one most likely to cause a production rollback if neglected. AI agents operating on customer data are subject to existing privacy regulations, internal data governance policies, and emerging AI-specific regulatory frameworks.

Einstein Trust Layer: Your Security Foundation

Salesforce's Einstein Trust Layer provides five core security capabilities that make Agentforce enterprise-grade by default. Understanding these capabilities is essential for governance sign-off from your legal, compliance, and information security teams.

Zero Data Retention

Prompts sent to the underlying LLM are not retained by Salesforce or the LLM provider. Each inference call is stateless from a data persistence perspective.

PII Masking

The Trust Layer can automatically detect and mask personally identifiable information before it is sent to the LLM, replacing it with tokens that are substituted back in the response.

Prompt Injection Defence

Jailbreak attempt detection and mitigation is built into the Trust Layer, preventing adversarial users from overriding agent instructions.

Audit Logging

Every agent interaction is logged in Salesforce with a full audit trail including the user, timestamp, action taken, and data accessed.

Toxicity Filtering

Output is screened for harmful content before being presented to the user, with configurable sensitivity thresholds per use case.

PII Classification Framework

Before deploying any agent that handles customer data, your organisation must classify all Salesforce fields into one of three categories and configure the Einstein Trust Layer accordingly.

Category	Examples	Trust Layer Setting	Handling
Tier 1 - Public	Account Name, Industry, City	Standard	No masking required
Tier 2 - Internal	Email, Phone, Job Title	PII Masking	Mask before LLM send

Category	Examples	Trust Layer Setting	Handling
Tier 3 - Sensitive	NI/Tax ID, DOB, Health data	Block	Never send to LLM

Regulatory Landscape for AI Agents in ASEAN & ANZ

Organisations deploying Agentforce in Southeast Asia and Australia face a patchwork of privacy and AI-related regulations. Understanding which apply to your deployment is a prerequisite for governance sign-off.

Jurisdiction	Key Regulation	Agentforce Implication
Australia	Privacy Act 1988 (amended 2023)	Consent required for automated decision-making affecting individuals; agents must not make binding decisions without human review
Singapore	PDPA + Model AI Governance Framework	Transparency obligations: users must know they are interacting with AI; explainability required for consequential decisions
Malaysia	PDPA 2010 (amendment in progress)	Sensitive data categories (health, finances) require explicit consent; agent must not store or process without documented basis
Philippines	Data Privacy Act 2012	Data Processing Agreement required; NPC registration for significant personal data processors; breach notification within 72 hours
Indonesia	PDP Law 2022	Data localisation may apply for certain sectors; cross-border transfer restrictions must be reviewed
UAE / KSA	DIFC DP Law / PDPL	Similar to GDPR in structure; legitimate interest basis available; DPO appointment may be required

Approval Workflows & Audit Requirements

Any agent action that modifies data (creating cases, updating accounts, processing orders) should have a corresponding approval workflow or at minimum a human-in-the-loop confirmation step for high-value transactions. Define your approval thresholds before go-live.

The key principle is proportionality: the approval overhead should match the risk of the action. Retrieving case history requires no approval. Creating a new case requires audit logging but not pre-approval. Issuing a refund over \$500 should require a human confirmation step. Cancelling a contract should always be escalated to a human agent.

Action Type	Example	Approval Requirement
Read-only retrieval	Look up case status, account details	No approval — log in audit trail
Non-destructive write	Create case, log activity, update stage	No pre-approval — post-audit review
Medium-risk write	Process refund <\$500, update SLA priority	Confirmation prompt to user before executing
High-risk write	Issue refund >\$500, escalate to executive	Escalate to human agent — do not auto-execute
Destructive action	Cancel contract, delete record, revoke access	Hard block — agent must always refuse and escalate

Compliance Checklist Before Go-Live

Legal sign-off obtained | PII classification complete | DPA / PDPA review done | Audit logging enabled and tested | Trust Layer configuration documented | Incident response plan for agent failures in place | Staff trained on agent limitations.

Chapter 7: The Readiness Scorecard

Answer each question with a score of 0 (not in place), 1 (partially in place), or 2 (fully in place). Sum your scores across all 25 questions. Maximum score: 50.

Scoring Guide

0-25: Not Ready — Significant foundational work required before any pilot deployment. 26-40: Partially Ready — Targeted remediation in 1-2 pillars needed; pilot possible in 3-6 months. 41-50: Ready to Pilot — Begin agent design and deployment. Kovil AI recommends a structured 4-week pilot.

Pillar 1: Data Quality (Questions 1-5)

#	Question	Score (0/1/2)
1	Account and Contact records are >80% complete for fields used by the target agent	_____
2	Duplicate records are identified and a remediation plan is in place	_____
3	A unified customer profile (Data Cloud or equivalent) is available	_____
4	PII fields are identified and classified in the Salesforce data model	_____
5	A data quality monitoring process is in place and reviewed regularly	_____

Pillar 2: Org Architecture (Questions 6-10)

#	Question	Score (0/1/2)
6	A dedicated Integration User exists (or can be created) for agent runtime	_____
7	All external APIs required by the agent are accessible via Named Credentials	_____
8	API limit headroom is >30% of daily allocation during peak hours	_____
9	A Full Copy (or Partial Copy) sandbox is available for agent UAT	_____
10	Org Health Check score is >75 and no critical security alerts are active	_____

Pillar 3: Use Case Readiness (Questions 11-15)

#	Question	Score (0/1/2)
11	A specific, scoped use case has been selected (not "general AI assistant")	_____
12	The use case has been mapped to the 2x2 impact/complexity matrix	_____

#	Question	Score (0/1/2)
13	Success metrics (deflection rate, AHT, CSAT) are defined and measurable	_____
14	The target process is documented end-to-end including exception paths	_____
15	Executive sponsor is identified and committed to the pilot	_____

Pillar 4: Team Enablement (Questions 16-20)

#	Question	Score (0/1/2)
16	A senior Salesforce Admin/Architect is allocated to the project (>50% time)	_____
17	Prompt engineering capability exists internally or via a partner	_____
18	A Business Analyst or Process Owner is assigned and has >25% time allocated	_____
19	A UAT plan with adversarial test cases is drafted or committed to	_____
20	A training and change management plan for end users is in place	_____

Pillar 5: Governance & Compliance (Questions 21-25)

#	Question	Score (0/1/2)
21	Legal and compliance teams are engaged and aware of the Agentforce pilot	_____
22	PII masking configuration has been reviewed against applicable privacy laws	_____
23	Audit logging is enabled and the log retention policy meets regulatory requirements	_____
24	An escalation path (agent to human handoff) is designed and tested	_____
25	An agent incident response plan exists and is communicated to relevant teams	_____

Total Score: _____ / 50

Interpreting Your Score by Pillar

In addition to your total score, examine your score by pillar. A low total score driven by a single pillar weakness is a very different situation from a low total score spread evenly across all five pillars. The former can often be resolved in 2-3 targeted sprints; the latter requires a more comprehensive foundational programme.

Pillar Score (0-10)	Interpretation	Recommended Action
0-3	Critical gap — blocks all agent deployment	Immediate remediation sprint; no pilot until resolved

Pillar Score (0-10)	Interpretation	Recommended Action
4-6	Significant gap — limits agent capability	Parallel remediation alongside scoped pilot
7-8	Minor gap — manageable risk	Document as known risk; monitor closely in pilot
9-10	Production ready for this pillar	No action required; maintain standards

Scoring Tip

Complete this scorecard with your Salesforce team lead, a business stakeholder, and your IT/security representative in the same session. Scores assigned by technical teams alone consistently overestimate readiness in the governance and team enablement pillars.

Chapter 8: Your Next Steps

Based on your Readiness Scorecard result, we recommend one of three paths. Whichever path you are on, Kovil AI is here to accelerate your journey to production.

Path A: Score 0-25 — Foundation First

What This Means

Your organisation has significant gaps across multiple pillars. Attempting a pilot now will likely fail — not because of the technology, but because the foundations are not ready.

- Engage Kovil AI for a free 90-minute Org Audit to identify the highest-priority gaps
- Initiate a data quality remediation sprint (4-6 weeks) focused on the target use case data
- Configure your Integration User, Named Credentials, and sandbox environment
- Begin prompt engineering training for your internal team
- Set a target: revisit this scorecard in 60 days with a goal of reaching 30+

Path B: Score 26-40 — Targeted Remediation

What This Means

You have solid foundations in some pillars but critical gaps in 1-2 areas. A pilot is achievable within 3-6 months with focused effort on the identified gaps.

- Identify your lowest-scoring pillar and address it first
- Run a 2-week sprint to remediate the data quality or architecture gaps
- Select your first use case using the 2x2 matrix — aim for high impact, low complexity
- Engage Kovil AI for a fixed-scope 4-week pilot delivery
- Define your success metrics and secure executive sponsor commitment

Path C: Score 41-50 — Ready to Pilot

What This Means

Congratulations — your organisation is genuinely ready to deploy Agentforce in production. The question is no longer whether you can, but how fast you should move.

- Begin agent design with Kovil AI's 4-Phase Implementation Framework immediately

- Target go-live within 4-6 weeks for your first use case
- Plan your Expand phase (3-5 agents) before the pilot is complete
- Establish an AI Centre of Excellence internally to govern future deployments
- Measure and report on ROI at 30, 60, and 90 days post go-live

Get Your Free Agentforce Org Audit

Kovil AI offers a complimentary 90-minute Agentforce readiness review for qualifying organisations. Our certified Salesforce architects will assess your org health, data quality, and use case fit — and deliver a prioritised action plan within 48 hours. Book your audit at kovil.ai/agentforce or email hello@kovil.ai.