



EBOOK · 24 PAGES · FREE DOWNLOAD

The Agentforce Readiness Guide

Is your Salesforce org ready for Agentforce? This guide covers the five readiness pillars — data quality, org architecture, use case prioritisation, team enablement, and governance — with a self-assessment scorecard included.

Salesforce

Agentforce

Einstein AI

Data Cloud

Apex

Produced by **Kovil AI Engineering Team** · kovil.ai · June 2026

Contents

01	Introduction — The Agentforce Moment	3
02	Pillar 1 — Data Quality & Data Cloud Readiness	5
03	Pillar 2 — Org Architecture	8
04	Pillar 3 — Use Case Prioritisation	11
05	Pillar 4 — Team Enablement	14
06	Pillar 5 — Governance & Einstein Trust Layer	17
07	Self-Assessment Scorecard	20
08	Next Steps with Kovil AI	23

01 / Introduction

The Agentforce Moment

Agentforce is Salesforce's platform for building, deploying, and governing autonomous AI agents directly inside your Salesforce org. Unlike chatbots that answer questions, Agentforce agents take actions — creating cases, updating records, sending emails, escalating to humans, and completing multi-step business processes without a human in the loop.

Agentforce runs on the Atlas Reasoning Engine — Salesforce's proprietary AI reasoning layer that interprets instructions, queries data, selects tools, and executes actions. Combined with Einstein Trust Layer for data security and Einstein 1 Data Cloud for real-time customer data, Agentforce represents the most complete out-of-the-box enterprise AI agent platform available today.

Who this guide is for: Salesforce admins, architects, and CRM leaders evaluating whether their org is ready to deploy Agentforce in production. Complete the scorecard in Chapter 7 before booking your first scoping call.

The 5 Agentforce Readiness Pillars

01 Data Quality

Is your Salesforce data clean, complete, and structured correctly for agents to reason over?

02 Org Architecture

Is your Salesforce org configured with the right metadata, flows, and permission sets to support agents?

03 Use Case Prioritisation

Have you identified the highest-ROI agent use case and defined a clear success metric?

04 Team Enablement

Does your team have the Salesforce admin, developer, and AI skills to build and sustain agents?

05 Governance

Are Einstein Trust Layer, data masking, audit trails, and agent action policies configured correctly?

02 / Pillar 1

Data Quality & Data Cloud Readiness

Agentforce agents reason over Salesforce data. If your data is incomplete, inconsistent, or stale, your agents will produce incorrect or embarrassing outputs. Data quality is the most frequently underestimated readiness pillar — and the one that most often causes production agent failures.

CRM Data Completeness

Agents that handle cases, leads, or opportunities need complete records — missing contact information, empty description fields, or unclassified record types cause agents to stall or hallucinate. Run a field completion audit before building agents.

Data Duplication

Duplicate accounts, contacts, and leads create conflicting signals for agents. Agentforce agents do not deduplicate — they surface what is in the org. Resolve duplicates using Salesforce Duplicate Management before agent deployment.

Data Cloud Integration

Einstein 1 Data Cloud allows agents to reason over unified customer profiles that combine Salesforce CRM data with external sources (e-commerce, marketing, support). If you have Data Cloud licensed, agents grounded in Data Cloud are significantly more accurate than agents grounded in CRM records alone.

Record Type Consistency

Agentforce agents use record types to determine context. Inconsistent record type usage — cases categorised differently across teams, opportunities missing stages — reduces agent accuracy on routing and triage tasks.

Flow Data Quality

If agents trigger or are triggered by Salesforce Flows, those flows must handle data correctly. Agents cannot recover from Flow errors — a failed flow causes the agent action to fail silently.

03 / Pillar 2

Org Architecture

Agentforce agents run inside your Salesforce org and interact with org metadata — objects, fields, flows, permission sets, and connected apps. Org architecture readiness determines whether agents can be deployed without disrupting existing functionality.

Salesforce Edition

Agentforce requires Salesforce Enterprise Edition or above with the Agentforce add-on license. Check your edition and license type before beginning any build.

Permission Sets & Profiles

Agent service users require specific permission sets — Agentforce User, Einstein Agent, and any custom permissions for the actions they will take. Permission configuration errors are the second most common deployment blocker.

Object & Field Accessibility

Agents access objects and fields via the agent's service user profile. Every object and field the agent needs to read or write must be accessible to that profile. Run a field-level security audit against the agent's planned actions.

Flow & Apex Actions

Agentforce agents invoke actions — Flows, Apex classes, and Prompt Templates. All invocable actions used by agents must be correctly labelled, described, and testable independently. Poorly described actions cause the Atlas Reasoning Engine to select the wrong action.

Connected Apps & OAuth

If agents need to call external APIs (via MuleSoft or direct HTTP), Connected App and OAuth configuration must be in place. Agents cannot authenticate to external services without a properly configured named credential.

Sandbox vs Production

Build and test agents in a sandbox that mirrors production data and metadata. Agents behave differently in sandboxes with stale data — full sandbox is strongly recommended.

04 / Pillar 3

Use Case Prioritisation

The right first Agentforce use case is not the most interesting one — it is the one most likely to succeed and generate measurable ROI. Use the scoring matrix below to evaluate candidate use cases.

Dimension	High (3)	Medium (2)	Low (1)
Data readiness	Complete, clean CRM records	Mostly complete, needs audit	Incomplete or external data
Volume	>1,000 interactions/month	200–1,000/month	<200/month
Decision clarity	Clear rules, lookup-based answers	Semi-structured decisions	Open-ended judgement
Labour saved	>20 hrs/week per agent/team	5–20 hrs/week	<5 hrs/week
Risk level	Low-risk, reversible actions	Medium risk with human review	High risk or irreversible
Stakeholder buy-in	Executive sponsor in place	Department-level support	No clear owner

Score each candidate use case. 14+: strong pilot candidate. Under 10: defer.

Top Agentforce use cases by industry

Financial Services

Loan status enquiries, compliance document checklists, account opening triage, fraud case routing.

Healthcare

Appointment scheduling, referral triage, benefits verification, prior authorisation status.

Retail / E-commerce

Order status, returns processing, product recommendations, loyalty points enquiries.

Technology / SaaS

Support case triage, licence renewal alerts, onboarding task automation, escalation routing.

Professional Services

Proposal status, project milestone updates, invoice query handling, NPS follow-up.

05 / Pillar 4

Team Enablement

Agentforce is new enough that most Salesforce teams — even experienced ones — lack the specific skills to build production agents. The skill gap is not a failure; it is simply a product of recency. Understanding it clearly is the fastest path to filling it.

Skill	Needed For	How to Acquire
Salesforce Admin (ADM 201)	All Agentforce deployments	Salesforce Trailhead — free
Agentforce Builder	Agent configuration, topic/action setup	Agentforce Trailhead module
Prompt Template authoring	Agent instruction quality	Einstein Prompts Trailhead
Apex development	Custom invocable actions	Developer Trailhead paths
Flow (Screen & Auto-launched)	Agent action triggers	Flow Trailhead paths
Data Cloud basics	Data Cloud-grounded agents	Data Cloud Specialist superbadge
LLM prompt engineering	Atlas instruction quality	Anthropic / OpenAI documentation

06 / Pillar 5

Governance & Einstein Trust Layer

The Einstein Trust Layer is Salesforce's enterprise AI safety architecture — built into the Salesforce platform and active for all Agentforce deployments. It provides data masking, toxicity detection, grounding with company data, prompt injection resistance, and a full audit trail of every agent action.

Data Masking

Sensitive fields (SSN, credit card, health data) are masked before being sent to the LLM. Configure which fields are masked in Einstein Trust Layer settings. This is not optional in regulated industries.

Audit Trail

Every agent action, reasoning step, and data access is logged to the Salesforce Audit Trail. Ensure your audit log retention policy covers the required period for your compliance framework.

Agent Action Policies

Define which actions agents are allowed to take, and under what conditions. An agent that can send emails must have clear guardrails on when and to whom. Action policies are configured per agent topic.

Human Escalation

Every production Agentforce agent must have a defined human escalation path. Configure escalation triggers — uncertainty threshold, user request, failed action — and test the escalation flow before go-live.

Toxicity & Content Safety

Einstein Trust Layer includes content safety filters that block toxic, harmful, or off-topic outputs. Review the filter configuration to ensure it matches your tolerance for false positives (legitimate content flagged) vs false negatives (harmful content passed through).

07 / Scorecard

Self-Assessment Scorecard

For each statement tick: ✓ Fully in place / ~ Partial / ✗ Not yet

Statement	✓	~	✗
Data Quality			
We have run a field completion audit on the key objects agents will use.	■	■	■
Duplicate records are resolved or suppressed in our Salesforce org.	■	■	■
Record types are used consistently across teams.	■	■	■
Data Cloud is configured and CRM data is unified (if licensed).	■	■	■
Org Architecture			
We have Agentforce-compatible Salesforce edition and licenses.	■	■	■
Service user permission sets for agents are configured and tested.	■	■	■
All objects/fields agents need are accessible to the agent service user.	■	■	■
Invocable actions (Flows, Apex) agents will use are correctly labelled.	■	■	■
Use Case			
We have scored candidate use cases using the framework in Chapter 4.	■	■	■
Our pilot use case scores 14+ on the scoring matrix.	■	■	■
A measurable success metric is defined for the pilot.	■	■	■
Executive sponsor is identified and briefed.	■	■	■
Governance			
Einstein Trust Layer data masking is configured for sensitive fields.	■	■	■
Audit log retention is set to meet our compliance requirements.	■	■	■
Human escalation triggers are defined for all planned agents.	■	■	■
Agent action policies are documented and reviewed by legal/compliance.	■	■	■

Score: 14–16 ✓ marks = Ready to build. 8–13 = Address gaps first. Under 8 = Engage a partner.

08 / Next Steps

Working with Kovil AI

Kovil AI is a specialist AI engineering firm. We build production Agentforce deployments for enterprise Salesforce orgs — fixed-price, outcome-based.

Agentforce Strategy & Readiness

\$8,000 fixed

Salesforce org audit, use case scoring, architecture blueprint. 2-week engagement.

Pilot Agent Build

\$15,000–\$28,000

One production Agentforce agent, one use case, fully deployed in your Salesforce org. 2–3 weeks.

Enterprise Agentforce Programme

\$45,000–\$120,000

Multiple agents, Data Cloud integration, MuleSoft connectivity, governance framework. 6–10 weeks.

Book a readiness call: 45 minutes with a Kovil AI Agentforce engineer. We review your scorecard, identify your top 3 gaps, and recommend a path to your first production Agentforce agent. kovil.ai · info@kovil.ai